

# LSAP HW4 Report

## 1 Investigating Modern Internet Infrastructure

### (1) Domain Analysis

本研究針對十個常用網站，蒐集其 DNS 基本紀錄：A (IPv4)、AAAA (IPv6)、CNAME (正規名稱)、MX (郵件交換)。另以 `dig +dnssec` 檢查回應旗標是否包含 ad (Authenticated Data)，據此判斷使用之遞迴解析器是否已完成 DNSSEC 簽章驗證。所有查詢統一透過公開遞迴解析器（例如 Cloudflare 1.1.1.1）執行，以降低本機快取與 ISP/地區策略差異的影響；該解析器支援 DNSSEC 驗證，便於以 ad 作為驗簽指標。

```
set -euo pipefail
DOMAINS_FILE="${1:-domains.txt}"
OUT_CSV="data/dns_records.csv"
TRACE_DIR="data/trace"
RESOLVER="${RESOLVER:-1.1.1.1}" # 可改 8.8.8.8
mkdir -p "$(dirname "$OUT_CSV")" "$TRACE_DIR"
echo "domain,ipv4,ipv6,cname,mx,dnssec_ad" > "$OUT_CSV"
join_semicolon() { sed 's/\$/ /' | paste -sd ';' -; }
has_ad() {
    local d="$1"
    dig @"$RESOLVER" +dnssec "$d" A +cmd +nocmd +noall 2>/dev/null \
    | awk '/flags:/ { if ($0 ~ / ad[ ;]/) {print 1} else {print 0}; exit }'
}
while IFS= read -r domain; do
    [[ -z "${domain// }" || "$domain" =~ ^# ]] && continue
    ipv4=$(dig @"$RESOLVER" +short A "$domain" | join_semicolon || true)
    ipv6=$(dig @"$RESOLVER" +short AAAA "$domain" | join_semicolon || true)
    cname=$(dig @"$RESOLVER" +short CNAME "$domain" | join_semicolon || true)
    mx=$(dig @"$RESOLVER" +short MX "$domain" | sed 's/\$/ /' | paste -sd ';' - || true)
    dnssec_ad=$((has_ad "$domain" || echo 0))
    dig +trace "$domain" > "$TRACE_DIR/$domain.trace.txt" 2>/dev/null || true
    echo "$domain,${ipv4:-},${ipv6:-},${cname:-},${mx:-},$dnssec_ad" >> "$OUT_CSV"
done < "$DOMAINS_FILE"
echo "Done → $OUT_CSV & $TRACE_DIR/*.trace.txt"
```

domain	ipv4	ipv6	cname	mx	dnssec_ad
scholar.google.com	scholar.google.com:142.250.66.68	scholar.google.com	scholar.google.com		
github.com	20.27.177.113		1.aspmx.google.com:10	10@1.aspmx.google.com:10@1.aspmx.google.com:5@2.aspmx.google.com	
stackoverflow.com	172.64.156.249.104.18.32.7		1.aspmx.google.com:10@2.aspmx.google.com:10@1.aspmx.google.com:5@1.aspmx.google.com		
colab.research.google.com	142.250.198.206	2404:6800:4012:9::20e	colab-av.research.google.com	10@3.aspmx.google.com:10@4.aspmx.google.com:5@2.aspmx.google.com	
kaggle.com	35.244.233.98		10@3.aspmx.google.com:10@4.aspmx.google.com:5@1.aspmx.google.com:5@2.aspmx.google.com		
overleaf.com	34.120.52.64		1.aspmx.google.com:10@3.aspmx.google.com:5@1.aspmx.google.com:5@2.aspmx.google.com		
regex101.com	78.47.220.195	2a01:4f8:1:1c1:1a4:-1	1.smtp.google.com		
diagrams.net	104.18.4.247;104.18.5.247	2606:4700:8812:47:2606:4700:8812:5ff	10@1.smtp.messagingengine.com:20		
codepen.io	104.18.147.32;104.18.163.32	2606:4700:8810:9300:2606:4700:8810:a30	1.aspmx.google.com:5@1.aspmx.google.com:5@2.aspmx.google.com:10@3.aspmx.google.com		

Figure 1: 10 個網站 DNS 紀錄彙整

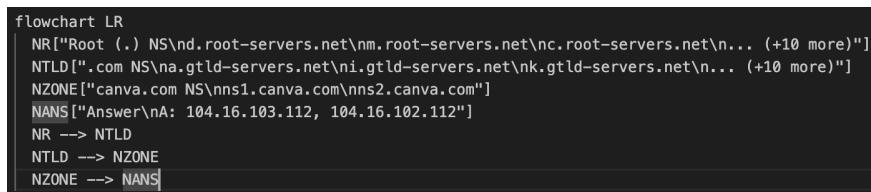


Figure 2: visual diagram of canva DNS lookup path

## (2) DNS Resolution Time Measurement

我們用 dig +stats 向指定遞迴解析器（預設 1.1.1.1）對 domains.txt 的 10 個網域各做多次查詢（預設 7 次），從輸出的「Query time」擷取每次解析耗時（毫秒），記錄 A（或設定 RECORD\_TYPE=AAAA）。

```
set -euo pipefail
DOMAINS_FILE="${1:-domains.txt}"
OUT_RAW="data/dns_time_raw.csv"
OUT_SUM="data/dns_time_avg.csv"
RESOLVER="${(RESOLVER:-1.1.1.1)}"
RECORD_TYPE="${(RECORD_TYPE:-A)}"
TRIALS="${(TRIALS:-7)}"
mkdir -p data
echo "domain,trial,ms" > "$OUT_RAW"
while IFS= read -r domain; do
    [[ -z "${domain// }" || "$domain" =~ ^# ]] && continue
    for t in $(seq 1 "$TRIALS"); do
        out=$(dig @${RESOLVER} +tries=1 +time=2 "$domain" "$RECORD_TYPE" +noall +answer +
            stats 2>/dev/null || true)
        ms=$((printf '%s\n' "$out" | awk '{print $2}' | awk '{print $1
            }')""
        [[ -z "$ms" ]] && ms=-1
        echo "$domain,$t,$ms" >> "$OUT_RAW"
        sleep "0.$(( RANDOM % 4 ))"
    done
done < "$DOMAINS_FILE"
awk -F, 'BEGIN{
    OFS=","; print "domain,trials,avg_ms,min_ms,max_ms,std_ms"
}
NR>1 {
    if ($3 >= 0) {
        n[$1]++; s[$1]+=$3; ss[$1]+=$3*$3;
        if (!( $1 in min ) || $3 < min[$1]) min[$1]=$3;
        if !( $1 in max ) || $3 > max[$1]) max[$1]=$3;
    }
}
END{
    for (d in n) {
        avg = s[d]/n[d];
        var = (ss[d]/n[d]) - (avg*avg); if (var < 0) var = 0;
        std = sqrt(var);
        print d, n[d], avg, min[d], max[d], std;
    }
}' "$OUT_RAW" | sort -t, -k3,3n > "$OUT_SUM"
echo "Done:"
echo " - $OUT_RAW"
echo " - $OUT_SUM"
```

domain	trials	avg_ms	min_ms	max_ms	std_ms
codepen.io	7	16.8571	7	34	8.57619
diagrams.net	7	22.2857	9	62	17.144
scholar.google.com	7	23.8571	7	96	29.763
colab.research.google.com	7	25.8571	9	94	28.0277
stackoverflow.com	7	33	7	138	44.4908
overleaf.com	7	40.2857	8	168	54.2763
kaggle.com	6	42.8333	9	141	47.397
regex101.com	7	43.5714	9	136	44.1454
github.com	7	46.8571	6	174	55.8632

Figure 3: DNS resolution times 之平均

## (3) DNS Load Balancing

對各網域連續查詢多次（A/AAAA），觀察是否出現不同的 IP 集合（unique\_answers>1），或僅順序不同（order\_varied=yes，常見 Round-Robin），以偵測 DNS 層級的負載平衡。

```

#!/usr/bin/env bash
set -euo pipefail
DOMAINS_FILE="${DOMAINS_FILE:-domains.txt}"
OUT_RAW="data/dns_lb_raw.csv"
OUT_SUM="data/dns_lb_summary.csv"
RESOLVER="${RESOLVER:-1.1.1.1}"
RECORD_TYPE="${RECORD_TYPE:-A}"
TRIALS="${TRIALS:-20}"
SUBNET_OPT="${SUBNET:++subnet=${SUBNET}}"
mkdir -p data
echo "domain,trial,ips_ordered" > "$OUT_RAW"
run_for_domain() {
    d="$1"
    [ -z "${d// }" ] && return
    case "$d" in
        *-* ) return;; esac
    for t in $(seq 1 "$TRIALS"); do
        ans=$(dig @${RESOLVER} ${SUBNET_OPT} +tries=1 +time=2 "$d" "$RECORD_TYPE" +noall +
              answer 2>/dev/null || true)
        ips=$(printf "%s\n" "$ans" | awk -v rt="$RECORD_TYPE" '$4==rt {print $5}' | sed 's/\.
        $// | paste -sd ';' -)'
        [ -z "$ips" ] && ips="(no-answer)"
        echo "$d,$t,$ips" >> "$OUT_RAW"
        sleep "0.$((RANDOM%3))"
    done
}
if [ -n "${DOMAINS:-}" ]; then
    set -- $DOMAINS
    for d in "$@"; do run_for_domain "$d"; done
else
    while IFS= read -r d; do run_for_domain "$d"; done < "$DOMAINS_FILE"
fi
awk -F, -v RT="$RECORD_TYPE" '
BEGIN{OFS=","; print "domain,trials,record_type,unique_ip_count,unique_sequences,
changed_in_pct,all_ips"}
NR>1{
    dom=$1; seq=$3; cnt[dom]++
    key=dom SUBSEP seq
    if(!(key in seqseen)) seqseen[key]=1
    n=split(seq,a,/;/)
    for(i=1;i<n;i++){ ip=a[i]; if(ip!="" && ip!="(no-answer)") ipseen[dom SUBSEP ip]=1 }
    if(!(dom in first)) first[dom]=seq
    if(seq!=first[dom]) chg[dom]++
}
END{
    for(k in ipseen){ split(k,p,SUBSEP); d=p[1]; ip=p[2]; if(!(d in uipcount)) uipcount[d]
        ]=0; if(!( (d SUBSEP ip) in touched)){ touched[d SUBSEP ip]=1; uiplist[d]=(uiplist[
        d] ";" ip); uipcount[d]++ } }
    for(k in seqseen){ split(k,p,SUBSEP); d=p[1]; useq[d]++ }
    for(d in cnt){
        ips = (d in uiplist)? substr(uiplist[d],2) : ""
        u = (d in uipcount)? uipcount[d] : 0
        us = (d in useq)? useq[d] : 0
        p = (chg[d]+0) * 100.0 / cnt[d]
        print d, cnt[d], RT, u, us, p, ips
    }
}
'

```

domain	trials	record_type	unique_ip_count	unique_sequences	changed_in_pct	all_ips
codepen.io	20	A	2	2	45	104.16.163.32;104.16.147.32
colab.research.google.com	18	A	4	15	88.8889	216.239.36.180;216.239.34.180;216.239.32.180;216.239.38.180
diagrams.net	20	A	2	2	50	104.18.5.247;104.18.4.247
github.com	20	A	1	1	0	20.27.177.113
kaggle.com	20	A	1	1	0	35.244.233.98
overleaf.com	20	A	1	1	0	34.120.52.64
regex101.com	20	A	1	1	0	78.47.220.195
scholar.google.com	20	A	1	1	0	142.250.66.68
stackoverflow.com	20	A	2	2	65	104.18.32.7;172.64.155.249

Figure 4: different IP responses

#### (4) CDN Identification

判斷各網站是否位於 CDN 後方，並辨識供應商（Cloudflare、Akamai、Fastly、CloudFront 等）；同時蒐集可佐證之「邊緣節點(Edge POP)」代碼(例如 CF-Ray 的 -TPE、CloudFront 的 X-Amz-Cf-Pop)。

```
chmod +x scripts/cdn_detect.sh
bash scripts/cdn_detect.sh
```

domain	variants	record	unique_ips	orgs	whois_countries	cdn_guess
colab.research.google.com	apex;www	A/AAAA	5 (A)			Unknown
diagrams.net	apex;www	A/AAAA	2 (A)	Cloudflare	Inc. (CLOUD14)	US
kaggle.com	apex;www	A/AAAA	1 (A)	Google LLC (GOOGL-2)	US	Google
github.com	apex;www	A/AAAA	2 (A)	Microsoft Corporation (MSFT)	US	Azure_or_Microsoft
scholar.google.com	apex;www	A/AAAA	3 (A)	Google LLC (GOOGL)	US	Google
regex101.com	apex;www	A/AAAA	1 (A)	CLOUD-NBG1	DE	Unknown
stackoverflow.com	apex;www	A/AAAA	2 (A)	Cloudflare	Inc. (CLOUD14)	US
codepen.io	apex;www	A/AAAA	2 (A)	Cloudflare	Inc. (CLOUD14)	US
overleaf.com	apex;www	A/AAAA	2 (A)	Google LLC (GOOGL-2)	US	Google

Figure 5: CDN provider and edge server locations

## (5) Network Performance Monitoring

針對十個網站量測三項網路效能指標：**平均往返延遲 (RTT, ms)**、**封包遺失率 (%)**、**下載吞吐 (Mbps)**，並彙整成表格。

domain	trials_ok	avg_rtt_ms	std_rtt_ms	avg_loss_pct	avg_speed_Mbps	std_speed_Mbps
codepen.io	5	8.3602	2.0363	0	0.567746	0.177329
colab.research.google.com	5	27.5198	16.5885	0	0.832413	0.132353
diagrams.net	5	21.1112	21.7087	0	0.0249312	0.0041775
github.com	5	92.4618	43.0267	0	1.18278	0.124296
kaggle.com	5	146.111	61.1482	4	0	0
overleaf.com	5	78.5538	64.8335	0	0.0050864	0.000204487
regex101.com	5	307.746	28.7614	0	0.285141	0.0273055
scholar.google.com	5	38.3042	16.3535	0	0.526838	0.146133
stackoverflow.com	5	69.808	61.2167	0	0	0

Figure 6: Network Performance

```

#!/usr/bin/env bash
set -euo pipefail
DOMAINS_FILE="${DOMAINS_FILE:-domains.txt}"
OUT_RAW="data/netperf_raw.csv"
OUT_SUM="data/netperf_summary.csv"
TRIALS="${TRIALS:-5}"
PING_COUNT="${PING_COUNT:-5}"
CURL_TIMEOUT="${CURL_TIMEOUT:-8}"
echo "domain,trial,rtt_ms,loss_pct,speed_Bps" > "$OUT_RAW"
run_one() {
    d="$1"
    [ -z "${d// }" ] && return
    case "$d" in
        *"/") return;; esac
    for t in $(seq 1 "$TRIALS"); do
        p=$(ping -c ${PING_COUNT} "$d" 2>&1 || true)
        loss=$(printf "%s\n" "$p" | grep -Eo '[0-9.]+% packet loss' | head -1 | sed 's/[^.]*\.\(.*\)\$/\1')
        [ -z "$loss" ] && loss=-1
        rtt=$(printf "%s\n" "$p" | grep -E 'round-trip|rtt' | awk -F=' ' '{print $2}' | awk '{print $1}' | awk -F'/' '{print $2}')
        [ -z "$rtt" ] && rtt=-1
        sp=$(curl -m ${CURL_TIMEOUT} -s -o /dev/null -w "%{speed_download}" "https://$d/" || true)
        [ -z "$sp" ] && sp=-1
        echo "$d,$t,$rtt,$loss,$sp" >> "$OUT_RAW"
        sleep "0.$((RANDOM%3))"
    done
}
while IFS= read -r d; do run_one "$d"; done < "$DOMAINS_FILE"
awk -F, '
function add(a,x){if(x>0){a["n"]++;a["s"]+=x;a["ss"]+=x*x}}
function avg(a){return a["n"]?a["s"]/a["n"]:-1}
function std(a){return a["n"]?sqrt(a["ss"]/a["n"]- (a["s"]/a["n"])^2):-1}
BEGIN{
    OFS=","
    print "domain","trials_ok","avg_rtt_ms","std_rtt_ms","avg_loss_pct","avg_speed_Mbps",
          "std_speed_Mbps"
}
NR>1{
    d=$1; rtt=$3+0; loss=$4+0; sp=$5+0
    if(rtt>0){rt[d,"n"]++; add(rt[d], rtt)}
    if(loss>=0){ls[d,"n"]++; add(ls[d], loss)}
    if(sp>0){sd[d,"n"]++; add(sd[d], sp)}
    seen[d]=1
}
END{
    for(d in seen){
        rta=avg(rt[d]); rts=std(rt[d])
        lsa=avg(ls[d])
        sda=avg(sd[d]); sds=std(sd[d])
        mbps=(sda>0)? sda*8/1000000 : -1
        mbpss=(sds>0)? sds*8/1000000 : -1
        n_ok=(rt[d,"n"]>sd[d,"n"])? rt[d,"n"] : sd[d,"n"]
        if(lsa>0 && lsa < 0) lsa=0
        print d, n_ok, rta, rts, lsa, mbps, mbpss)
    }' "$OUT_RAW" | sort -t, -k1,1 > "$OUT_SUM"
echo "Done:"
echo " - $OUT_RAW"
echo " - $OUT_SUM"

```

## (6) Network Routing Path Analysis

traceroute 量測從本機到 cloudflare.com 的實際路由，擷取每一跳的回應時間；再用自動化腳本將輸出解析為表格，計算各 hop 的平均往返延遲 (ms)。同時對每個 IP 做反向 DNS 與 IP 資訊查詢，補齊 \*\*Hostname / Organization (ISP) / Country / City / 經緯度 \*\*。

```

#!/usr/bin/env bash
set -euo pipefail
D="${1:?usage: route_analyze.sh <domain>}"
OUT_CSV="data/route_${D}.csv"
OUT_MMD="report/route_${D}.mmd"
TR="$(command -v traceroute || true)"
[ -z "$TR" ] && { echo "traceroute not found"; exit 1; }
echo "hop,ip,hostname,org,country,location,avg_ms" > "$OUT_CSV"
$TR -n -q 3 -w 2 "$D" | awk 'NR>1' | while read -r line; do
    hop=$(echo "$line" | awk "{print \$1}")
    ip=$(echo "$line" | grep -Eo '([0-9]{1,3}\.){3}[0-9]{1,3}' | head -1)

```

```

if [ -z "$ip" ]; then
    echo "$line" | grep -q '\* \* \* && ip="*"
fi
if [ "$ip" = "*" ] || [ -z "$ip" ]; then
    hn="*"; org=""; ctry=""; loc=""; avg="-1"
else
    hn=$(dig +short -x "$ip" | sed 's/\.\$//' | head -1)
    [ -z "$hn" ] && hn="-"
    ms_raw=$(echo "$line" | grep -Eo '[0-9]+\.[0-9]+ ms' | awk '{print $1}')
    if [ -z "$ms_raw" ]; then avg="-1"; else
        c=0; s=0; echo "$ms_raw" | while read -r v; do c=$((c+1)); s=$(awk -v a="$s" -v b="$
            $v" 'BEGIN{printf "%."6f", a+b}'); done
        avg=$(awk -v s="$s" -v c="$c" 'BEGIN{ if(c>0) printf "%.3f", s/c; else print "-1"
            }')"
    fi
    wf=$(whois "$ip" 2>/dev/null || true)"
    org=$(printf "%s\n" "$wf" | awk -F': *' 'tolower($1)~/^(orgname|org-name|owner|
        organization|descr|netname)$/{print $2; exit}'")
    ctry=$(printf "%s\n" "$wf" | awk -F': *' 'tolower($1)~/^country$/ {print $2; exit}'")"
    [ -z "$org" ] && org="-"; [ -z "$ctry" ] && ctry="-"
    ji=$(curl -m 3 -s "https://ipinfo.io/$ip" || true)"
    city=$(printf "%s" "$ji" | sed -n 's/.*"city"[:space:]*:[[:space:]]*"\([^\"]*\)\"
        .*/\1/p' | head -1)"
    region=$(printf "%s" "$ji" | sed -n 's/.*"region"[:space:]*:[[:space:]]*"\([^\"]*\)\"
        .*/\1/p' | head -1)"
    if [ -n "$city$region" ]; then loc="$city/$region"; else loc="-"; fi
fi
echo "$hop,$ip,$hn,$org,$ctry,$loc,$avg" >> "$OUT_CSV"
done
echo "flowchart LR" > "$OUT_MMD"
echo " classDef hop fill:#eef,stroke:#999,rx:10,ry:10;" >> "$OUT_MMD"
i=0
while IFS=, read -r hop ip hn org ctry loc avg; do
    [ "$hop" = "hop" ] && continue
    label="Hop ${hop}\n${ip}\n${hn}\n${org}\n${ctry}\n${loc}\n${avg}\n${ms}"
    echo " N${hop}[\"$label\"]:::hop" >> "$OUT_MMD"
    if [ "$i" -gt 0 ]; then prev=$((hop-1)); echo " N${prev} --> N${hop}" >> "$OUT_MMD"; fi
    i=$((i+1))
done < "$OUT_CSV"
echo "Done"
echo "$OUT_CSV"
echo "$OUT_MMD"

```

route_kaggle.com						
hop	ip	hostname	org	country	location	avg_ms
2	123.51.152.254	-	NCICNET-NET	TW	Taipei/Taiwan	8.700
3	220.228.20.157	-	NCICNET-TW	TW	Taipei/Taiwan	5.799
220.228.20.149	220.228.20.149	-	NCICNET-TW	TW	Taipei/Taiwan	9.880
4	192.72.107.189	h189-192-72-107.seed.net.tw	APNIC-ERX-192-72-3-0	AU	Taipei/Taiwan	12.735
192.72.107.97	192.72.107.97	h97-192-72-107.seed.net.tw	APNIC-ERX-192-72-3-0	AU	Taipei/Taiwan	6.585
192.72.107.189	192.72.107.189	h189-192-72-107.seed.net.tw	APNIC-ERX-192-72-3-0	AU	Taipei/Taiwan	5.589
5	139.175.56.141	r56-141.seed.net.tw	APNIC-ERX-139-175-0-0	AU	Taipei/Taiwan	7.672
139.175.57.185	139.175.57.185	r57-185.seed.net.tw	APNIC-ERX-139-175-0-0	AU	Taipei/Taiwan	6.026
6	139.175.58.202	r58-202.seed.net.tw	APNIC-ERX-139-175-0-0	AU	Taipei/Taiwan	8.155
139.175.59.163	139.175.59.163	r59-163.seed.net.tw	APNIC-ERX-139-175-0-0	AU	Taipei/Taiwan	7.132
139.175.59.167	139.175.59.167	r59-167.seed.net.tw	APNIC-ERX-139-175-0-0	AU	Taipei/Taiwan	5.625
7	139.175.59.217	r59-217.seed.net.tw	APNIC-ERX-139-175-0-0	AU	Taipei/Taiwan	7.201
8	142.250.172.86	-	GOOGLE	US	Taipei/Taiwan	5.718
35.244.233.98	35.244.233.98	98.233.244.35.bc.googleusercontent.com	GOOGLE-CLOUD	US	Kansas City/Missouri	5.899
142.250.172.86	142.250.172.86	-	GOOGLE	US	Taipei/Taiwan	5.771

Figure 7: kaggle's traceroute record



Figure 8: kaggls's route diagram (節錄)

## (7) Backend Server Investigation

我們以 curl -I -L 摷取各網站最終回應標頭 (follow redirects)，從 Server、X-Powered-By、Via 與常見雲端／CDN 指紋 (如 CF-Ray, CloudFront, Fastly, gws, ESF 等) 判斷背後的 Web 伺服器與雲邊緣層。流程自動化：逐一對清單中的網域發請求，解析標頭欄位並彙整為 CSV，比對「伺服器核心 (Nginx/Apache/LiteSpeed/OpenResty 等)」與「CDN／雲端代理 (Cloudflare、Google Frontend、Akamai、CloudFront、Vercel 等)」兩個層面。

```

#!/usr/bin/env bash
set -euo pipefail
DOMAINS_FILE="${DOMAINS_FILE:-domains.txt}"
OUT_RAW="data/backend_raw.csv"
OUT_SUM="data/backend_summary.csv"
TIMEOUT="${TIMEOUT:-8}"
UA="${UA:-Mozilla/5.0}"
SCHEMES="${SCHEMES:-https http}"
VARIANTS="${VARIANTS:-apex www}"
echo "domain,variant,scheme,status,server,via,x_powered_by,cdn_hint,app_hint" > "$OUT_RAW"
"

fetch(){
    u="$1"; h=$(curl -m "$TIMEOUT" -A "$UA" -sS -I -L -k "$u" -D - -o /dev/null || true)
    st=$(printf "%s" "$h" | awk 'BEGIN{RS="\r?\n\r?\n"} NR==1{print}' | awk 'toupper($0) ~
        ^HTTP/{code=$2} END{print code+0}')
    sv=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)=="server"{print
        $2; exit}')
    via=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)=="via"{print $2;
        exit}')
    xp=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)=="x-powered-by"{
        print $2; exit}')
    cf=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)=="cf-ray"{print "
        Cloudflare"; exit}')
    ak=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)~/^server-timing${
        /{if(tolower($2)~/^(ak_)/) print "Akamai"}')
    fa=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)=="x-served-by"{
        (tolower($2)~/^(fastly)/) print "Fastly"}')
    ec=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)~/^(x-ec-)/{print
        "Edgio_Limelight"; exit}')
    lb=$(printf "%s" "$h" | awk -F': *' 'BEGIN{IGNORECASE=1} tolower($1)=="server"{if(
        tolower($2)~/^(gws|tsa)/) print "Google"; else if(tolower($2)~/^(cloudfront)/) print
        "AWS_CloudFront"; else if(tolower($2)~/^(azure)/) print "Azure"}')
    ch="${cf:-}${ak:+${ak}}${fa:+${fa}}${ec:+${ec}}${lb:+${lb}}"; ch=$(printf "%s" "$ch" | sed
        's/^;*//;s/;*$/'; tr ';' '\n' | sort -u | paste -sd ';' -)
    app=$(printf "%s" "$sv;$xp" | tr '[:upper:]' '[:lower:]')
    app=$(printf "%s" "$app" | awk -F';' '{for(i=1;i<NF;i++) s=$i; if(s~/nginx/) a="nginx"
        ; if(s~/apache/) a=(a?a;"")"apache"; if(s~/litespeed/) a=(a?a;"")"litespeed";
        if(s~/iis/) a=(a?a;"")"microsoft_iis"; if(s~/gws/) a=(a?a;"")"google_gws";
        if(s~/tsa/) a=(a?a;"")"google_tsa"; if(s~/cloudflare/) a=(a?a;"")"cloudflare"
    ')
}

```

```

; if(s~/cloudfront/) a=(a?a";":")"aws_cloudfront"; if(s~/ats|apache traffic server
/) a=(a?a";":")"apache_traffic_server"}; print a; exit')
echo "$st|$sv|$via|$xp|$ch|$app"
}
run_one(){
d="$1"; v="$2"; s="$3"; t="$d"; [ "$v" = "www" ] && t="www.$d"; u="$s://$t/"
r=$(fetch "$u"); st="${r%/*}"; rest="${r##*}"; sv="${rest%/*}"; rest="${rest##*}";
vi="${rest%/*}"; rest="${rest##*}"; xp="${rest%/*}"; rest="${rest##*}"; cdn="${rest%/*}";
app="${rest##*}"
echo "$d,$v,$s,${st:-},${sv:-},${vi:-},${xp:-},${cdn:-},${app:-}" >> "$OUT_RAW"
}
while IFS= read -r d; do
[ -z "${d// }" ] && continue
case "$d" in \#*) continue;; esac
for v in $VARIANTS; do for s in $SCHEMES; do run_one "$d" "$v" "$s"; done; done
done < "$DOMAINS_FILE"
awk -F, 'BEGIN{OFS=","; print "domain,server_guess,cdn_hint,examples"}'
NR>1{
k=$1
if($9!=""){split($9,a,/); for(i in a) app[k ":" a[i]]=1}
if($8!=""){split($8,c,/); for(i in c) cdn[k ":" c[i]]=1}
ex[k]=$2" "$3" "$4" "$5
seen[k]=1
}
END{
for(d in seen){
sg="" sep=""
for(x in app){ n=split(x,p,":"); if(p[1]==d){ sg=sg sep p[2]; sep=";" }
cg="" sep=""
for(x in cdn){ n=split(x,p,":"); if(p[1]==d){ cg=cg sep p[2]; sep=";" }
if(sg=="") sg="-"; if(cg=="") cg="-"
print d,sg,cg,ex[d]
}
}
"$OUT_RAW" | sort -t, -k1,1 > "$OUT_SUM"
echo "Done:"
echo " - $OUT_RAW"
echo " - $OUT_SUM"

```

domain	server_guess	cdn_hint	examples
<a href="http://codepen.io">codepen.io</a>	cloudflare	Cloudflare	www http 301 cloudflare
<a href="http://colab.research.google.com">colab.research.google.com</a>	-	-	www http 0
<a href="http://diagrams.net">diagrams.net</a>	cloudflare	Cloudflare	www http 301 cloudflare
<a href="http://github.com">github.com</a>	-	-	www http 301 github.com
<a href="http://kaggle.com">kaggle.com</a>	-	-	www http 404
<a href="http://overleaf.com">overleaf.com</a>	nginx	-	www http 308 nginx
<a href="http://regex101.com">regex101.com</a>	nginx	-	www http 301 nginx
<a href="http://scholar.google.com">scholar.google.com</a>	-	-	www http 302 scholar
<a href="http://stackoverflow.com">stackoverflow.com</a>	cloudflare	Cloudflare	www http 302 cloudflare

Figure 9: route diagram